

## [Hugging Face](#)

- [Models](#)
- [Datasets](#)
- [Spaces](#)
- [Posts](#)
- [Docs](#)
- [Pricing](#)



[BAAI](#)

/

[bge-large-zh-v1.5](#)

**like224**

[Feature Extractionsentence-transformersPyTorchSafetensorsTransformersChinesebertsentence-similarityInference](#)  
[Endpoints5 papers](#)  
License:[mit](#)

[Model card](#)[Files and versions](#)[Community](#)

[15](#)

Deploy

Use in libraries

[Edit model card](#)

# FlagEmbedding

[Model List](#) | [FAQ](#) | [Usage](#) | [Evaluation](#) | [Train](#) | [Contact](#) | [Citation](#) | [License](#)

For more details please refer to our Github: [FlagEmbedding](#).

If you are looking for a model that supports more languages, longer texts, and other retrieval methods, you can try using [bge-m3](#).

[English](#) | [中文](#)

FlagEmbedding focuses on retrieval-augmented LLMs, consisting of the following projects currently:

- **Long-Context LLM:** [Activation Beacon](#)
- **Fine-tuning of LM :** [LM-Cocktail](#)
- **Dense Retrieval:** [BGE-M3](#), [LLM Embedder](#), [BGE Embedding](#)
- **Reranker Model:** [BGE Reranker](#)
- **Benchmark:** [C-MTEB](#)

## News

- 1/30/2024: Release **BGE-M3**, a new member to BGE model series! M3 stands for **M**ulti-linguality (100+ languages), **M**ulti-granularities (input length up to 8192), **M**ulti-Functionality (unification of dense, lexical, multi-vec/colbert retrieval). It is the first embedding model which supports all three retrieval methods, achieving new SOTA on multi-lingual (MIRACL) and cross-lingual (MKQA) benchmarks. [Technical Report](#) and [Code](#). :fire:
- 1/9/2024: Release [Activation-Beacon](#), an effective, efficient, compatible, and low-cost (training) method to extend the context length of LLM. [Technical Report](#) :fire:
- 12/24/2023: Release **LLaRA**, a LLaMA-7B based dense retriever, leading to state-of-the-art performances on MS MARCO and BEIR. Model and code will be open-sourced. Please stay tuned. [Technical Report](#) :fire:
- 11/23/2023: Release [LM-Cocktail](#), a method to maintain general capabilities during fine-tuning by merging multiple language models. [Technical Report](#) :fire:
- 10/12/2023: Release [LLM-Embedder](#), a unified embedding model to support diverse retrieval augmentation needs for LLMs. [Technical Report](#)
- 09/15/2023: The [technical report](#) and [massive training data](#) of BGE has been released
- 09/12/2023: New models:
  - **New reranker model:** release cross-encoder models BAAI/bge-reranker-base and BAAI/bge-reranker-large, which are more powerful than embedding model. We recommend to use/fine-tune them to re-rank top-k documents returned by embedding models.

- **update embedding model:** release `bge-*`-v1.5 embedding model to alleviate the issue of the similarity distribution, and enhance its retrieval ability without instruction.

More

- 
- 
- 
- 
- 

## Model List

`bge` is short for BAAI general embedding.

Model	Language		Description	query instruction for retrieval [1]
<a href="#">BAAI/bge-m3</a>	Multilingual	<a href="#">Inference</a> <a href="#">Fine-tune</a>	Multi-Functionality(dense retrieval, sparse retrieval, multi-vector(colbert)), Multi-Linguality, and Multi-Granularity(8192 tokens)	
<a href="#">BAAI/llm-embedder</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a unified embedding model to support diverse retrieval augmentation needs for LLMs	See <a href="#">README</a>
<a href="#">BAAI/bge-reranker-large</a>	Chinese and English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a cross-encoder model which is more accurate but less efficient [2]	

Model	Language		Description	query instruction for retrieval [1]
<a href="#">BAAI/bge-reranker-base</a>	Chinese and English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a cross-encoder model which is more accurate but less efficient [2]	
<a href="#">BAAI/bge-large-en-v1.5</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	version 1.5 with more reasonable similarity distribution	Represent this sentence for searching relevant passages:
<a href="#">BAAI/bge-base-en-v1.5</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	version 1.5 with more reasonable similarity distribution	Represent this sentence for searching relevant passages:
<a href="#">BAAI/bge-small-en-v1.5</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	version 1.5 with more reasonable similarity distribution	Represent this sentence for searching relevant passages:

Model	Language		Description	query instruction for retrieval [1]
<a href="#">BAAI/bge-large-zh-v1.5</a>	Chinese	<a href="#">Inference</a> <a href="#">Fine-tune</a>	version 1.5 with more reasonable similarity distribution	<div> <div>为这个句子生成</div> <div>表示以用于检索</div> <div>相关文章：</div> </div>
<a href="#">BAAI/bge-base-zh-v1.5</a>	Chinese	<a href="#">Inference</a> <a href="#">Fine-tune</a>	version 1.5 with more reasonable similarity distribution	<div> <div>为这个句子生成</div> <div>表示以用于检索</div> <div>相关文章：</div> </div>
<a href="#">BAAI/bge-small-zh-v1.5</a>	Chinese	<a href="#">Inference</a> <a href="#">Fine-tune</a>	version 1.5 with more reasonable similarity distribution	<div> <div>为这个句子生成</div> <div>表示以用于检索</div> <div>相关文章：</div> </div>
<a href="#">BAAI/bge-large-en</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	rank 1st in <a href="#">MTEB</a> leaderboard	<div> <div>Represent this sentence for searching relevant passages:</div> </div>
<a href="#">BAAI/bge-base-en</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a base-scale model but with similar ability to bge-large-en	<div> <div>Represent this sentence for</div> </div>

Model	Language		Description	query instruction for retrieval [1]
				searching relevant passages:
<a href="#">BAAI/bge-small-en</a>	English	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a small-scale model but with competitive performance	Represent this sentence for searching relevant passages:
<a href="#">BAAI/bge-large-zh</a>	Chinese	<a href="#">Inference</a> <a href="#">Fine-tune</a>	:trophy: rank <b>1st</b> in <a href="#">C-MTEB</a> benchmark	为这个句子生成表示以用于检索相关文章：
<a href="#">BAAI/bge-base-zh</a>	Chinese	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a base-scale model but with similar ability to bge-large-zh	为这个句子生成表示以用于检索相关文章：
<a href="#">BAAI/bge-small-zh</a>	Chinese	<a href="#">Inference</a> <a href="#">Fine-tune</a>	a small-scale model but with competitive performance	为这个句子生成表示以用于检索相关文章：

[1]: If you need to search the relevant passages to a query, we suggest to add the instruction to the query; in other cases, no instruction is needed, just use the original query directly. In all cases, **no instruction** needs to be added to passages.

[2]: Different from embedding model, reranker uses question and document as input and directly output similarity instead of embedding. To balance the accuracy and time cost, cross-encoder is widely used to re-rank top-k documents retrieved by other simple models. For examples, use bge embedding model to retrieve top 100 relevant documents, and then use bge reranker to re-rank the top 100 document to get the final top-3 results.

All models have been uploaded to Huggingface Hub, and you can see them at <https://huggingface.co/BAAI>. If you cannot open the Huggingface Hub, you also can download the models at <https://model.baai.ac.cn/models>.

## Frequently asked questions

1. How to fine-tune bge embedding model?

- 
- 
- 

2. The similarity score between two dissimilar sentences is higher than 0.5

3. When does the query instruction need to be used

## Usage

### Usage for Embedding Model

Here are some examples for using bge models with [FlagEmbedding](#), [SentenceTransformers](#), [Langchain](#), or [Huggingface Transformers](#).

#### Using FlagEmbedding

```
pip install -U FlagEmbedding
```

If it doesn't work for you, you can see [FlagEmbedding](#) for more methods to install FlagEmbedding.

```
from FlagEmbedding import FlagModel
```

```
sentences_1 = ["样例数据-1", "样例数据-2"]
```

```
sentences_2 = ["样例数据-3", "样例数据-4"]
```

```

model = FlagModel('BAAI/bge-large-zh-v1.5',

query_instruction_for_retrieval="为这个句子生成表示以用于检索相关文章：",

use_fp16=True) # Setting use_fp16 to True speeds up computation with a slight performance
degradation

embeddings_1 = model.encode(sentences_1)

embeddings_2 = model.encode(sentences_2)

similarity = embeddings_1 @ embeddings_2.T

print(similarity)

# for s2p(short query to long passage) retrieval task, suggest to use encode_queries() which will
automatically add the instruction to each query

# corpus in retrieval task can still use encode() or encode_corpus(), since they don't need instruction

queries = ['query_1', 'query_2']

passages = ["样例文档-1", "样例文档-2"]

q_embeddings = model.encode_queries(queries)

p_embeddings = model.encode(passages)

scores = q_embeddings @ p_embeddings.T

```

For the value of the argument `query_instruction_for_retrieval`, see [Model List](#).

By default, FlagModel will use all available GPUs when encoding. Please set `os.environ["CUDA_VISIBLE_DEVICES"]` to select specific GPUs. You also can set `os.environ["CUDA_VISIBLE_DEVICES"]=""` to make all GPUs unavailable.

## Using Sentence-Transformers

You can also use the `bge` models with [sentence-transformers](#):



```
pip install -U sentence-transformers
```

```
from sentence_transformers import SentenceTransformer
```

```
sentences_1 = ["样例数据-1", "样例数据-2"]
```

```
sentences_2 = ["样例数据-3", "样例数据-4"]
```

```
model = SentenceTransformer('BAAI/bge-large-zh-v1.5')
```

```
embeddings_1 = model.encode(sentences_1, normalize_embeddings=True)
```

```
embeddings_2 = model.encode(sentences_2, normalize_embeddings=True)
```

```
similarity = embeddings_1 @ embeddings_2.T
```

```
print(similarity)
```

For s2p(short query to long passage) retrieval task, each short query should start with an instruction (instructions see [Model List](#)). But the instruction is not needed for passages.

```
from sentence_transformers import SentenceTransformer
```

```
queries = ['query_1', 'query_2']
```

```
passages = ["样例文档-1", "样例文档-2"]
```

```
instruction = "为这个句子生成表示以用于检索相关文章： "
```

```
model = SentenceTransformer('BAAI/bge-large-zh-v1.5')
```

```
q_embeddings = model.encode([instruction+q for q in queries], normalize_embeddings=True)
```

```
p_embeddings = model.encode(passages, normalize_embeddings=True)
```

```
scores = q_embeddings @ p_embeddings.T
```

## Using Langchain

You can use `bge` in langchain like this:

```
from langchain.embeddings import HuggingFaceBgeEmbeddings

model_name = "BAAI/bge-large-en-v1.5"

model_kwargs = {'device': 'cuda'}

encode_kwargs = {'normalize_embeddings': True} # set True to compute cosine similarity

model = HuggingFaceBgeEmbeddings(

    model_name=model_name,

    model_kwargs=model_kwargs,

    encode_kwargs=encode_kwargs,

    query_instruction="为这个句子生成表示以用于检索相关文章： "

)

model.query_instruction = "为这个句子生成表示以用于检索相关文章： "
```

## Using HuggingFace Transformers

With the transformers package, you can use the model like this: First, you pass your input through the transformer model, then you select the last hidden state of the first token (i.e., [CLS]) as the sentence embedding.

```
from transformers import AutoTokenizer, AutoModel

import torch

# Sentences we want sentence embeddings for

sentences = ["样例数据-1", "样例数据-2"]
```

```

# Load model from HuggingFace Hub

tokenizer = AutoTokenizer.from_pretrained('BAAI/bge-large-zh-v1.5')

model = AutoModel.from_pretrained('BAAI/bge-large-zh-v1.5')

model.eval()


# Tokenize sentences

encoded_input = tokenizer(sentences, padding=True, truncation=True, return_tensors='pt')

# for s2p(short query to long passage) retrieval task, add an instruction to query (not add instruction for passages)

# encoded_input = tokenizer([instruction + q for q in queries], padding=True, truncation=True, return_tensors='pt')


# Compute token embeddings

with torch.no_grad():

    model_output = model(**encoded_input)

    # Perform pooling. In this case, cls pooling.

    sentence_embeddings = model_output[0][:, 0]

    # normalize embeddings

    sentence_embeddings = torch.nn.functional.normalize(sentence_embeddings, p=2, dim=1)

    print("Sentence embeddings:", sentence_embeddings)

```

## Usage for Reranker

Different from embedding model, reranker uses question and document as input and directly output similarity instead of embedding. You can get a relevance score by inputting query and

passage to the reranker. The reranker is optimized based cross-entropy loss, so the relevance score is not bounded to a specific range.

### Using FlagEmbedding

```
pip install -U FlagEmbedding
```

Get relevance scores (higher scores indicate more relevance):

```
from FlagEmbedding import FlagReranker
```

```
reranker = FlagReranker('BAAI/bge-reranker-large', use_fp16=True) # Setting use_fp16 to True speeds up computation with a slight performance degradation
```

```
score = reranker.compute_score(['query', 'passage'])
```

```
print(score)
```

```
scores = reranker.compute_score(['what is panda?', 'hi'], ['what is panda?', 'The giant panda (Ailuropoda melanoleuca), sometimes called a panda bear or simply panda, is a bear species endemic to China.'])
```

```
print(scores)
```

### Using Huggingface transformers

```
import torch
```

```
from transformers import AutoModelForSequenceClassification, AutoTokenizer
```

```
tokenizer = AutoTokenizer.from_pretrained('BAAI/bge-reranker-large')
```

```
model = AutoModelForSequenceClassification.from_pretrained('BAAI/bge-reranker-large')
```

```
model.eval()
```

```
pairs = [['what is panda?', 'hi'], ['what is panda?', 'The giant panda (Ailuropoda melanoleuca), sometimes called a panda bear or simply panda, is a bear species endemic to China.']]
```

```
with torch.no_grad():
```

```
    inputs = tokenizer(pairs, padding=True, truncation=True, return_tensors='pt', max_length=512)
```

```
    scores = model(**inputs, return_dict=True).logits.view(-1, ).float()
```

```
    print(scores)
```